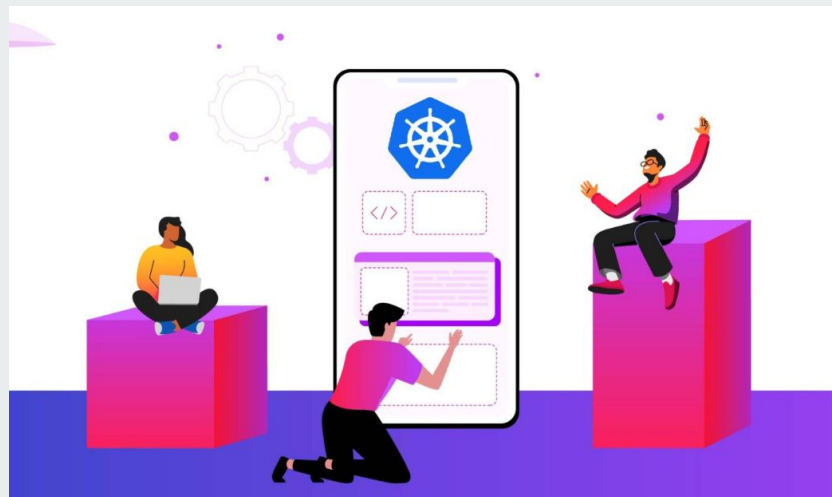


Kubernetes 5

Volume & StatefulSet

caodangsao@gmail.com



Module Target

Kết thúc bài học, học viên cần đạt được các kỹ năng sau:

- Có thể hiểu và tạo được Volume trong Kubernetes.
- Nắm được các đối tượng để khởi tạo Volume như StorageClass, PersistentVolume và PersistentVolumeClaim.
- Có khả năng triển khai ứng dụng dưới dạng Statefulsets.

Nội dung

1 Volume in K8s

- Quản lý cấu hình ứng dụng
- Volume là gì
- Volume Type

2 PV, PVC and Statefulset

- PV, PVC
- Storage Class
- Statefulset

3 Summary

- Tóm tắt nội dung buổi học



Lab 0: Manage Application Configuration



Lab 1: Sharing a volume, in action



Lab 2: Volume type: hostPath



Lab 3: Using PV, PVC



Lab 4: StatefulSets



Section 1:

Volume in K8s

Manage Application Configuration

- Cấu hình ứng dụng là gì ?
- ConfigMaps
- Secrets
- Environment Variables
- Configuration Volumes
- Lab demo

Application Configuration

- Ta có thể hiểu "**cấu hình**" là các cài đặt ảnh hưởng đến hoạt động của ứng dụng.
- Kubernetes cho phép người dùng chuyển các giá trị cấu hình động cho ứng dụng lúc Runtime.
- Các cấu hình động này giúp người dùng kiểm soát luồng ứng dụng.

ConfigMap

- Giữ dữ liệu không nhạy cảm trong ConfigMap, dữ liệu này có thể được chuyển đến container của ứng dụng.
- Config Map lưu trữ data ở định dạng Key-Value.
- ConfigMaps cho phép ta tách các cấu hình của mình khỏi Pod và các thành phần.
- ConfigMap giúp làm cho các cấu hình dễ thay đổi và quản lý hơn, đồng thời ngăn mã hóa cứng dữ liệu cấu hình thành các thông số kỹ thuật của Pod.

ConfigMap Commands

→ Tạo ConfigMaps qua File cấu hình.

```
kubectl create configmap [NAME] --from-file [/PATH/TO/FILE.PROPERTIES] --from-file  
[/PATH/TO/FILE2.PROPERTIES]
```

```
kubectl create configmap [NAME] --from-file [/PATH/TO/DIRECTORY]
```

→ Lấy ConfigMap qua CLI

```
kubectl get configmap <Config_map_name> -o yaml/json
```


Secrets

- **Secrets** tương tự như ConfigMap nhưng được thiết kế để giữ các dữ liệu nhạy cảm.
- Tạo Secrets từ file.

```
kubectl create secret generic db-user-pass --from-file=./username.txt --from-file=./password.txt
```

- Lưu ý: Các ký tự đặc biệt như `$`, `\`, `*` và `!` yêu cầu **escaping**.

Secrets

→ Lấy Secrets.

```
kubectl get secrets
```

→ Mô tả Secrets.

```
kubectl describe secrets <Secret_name>
```

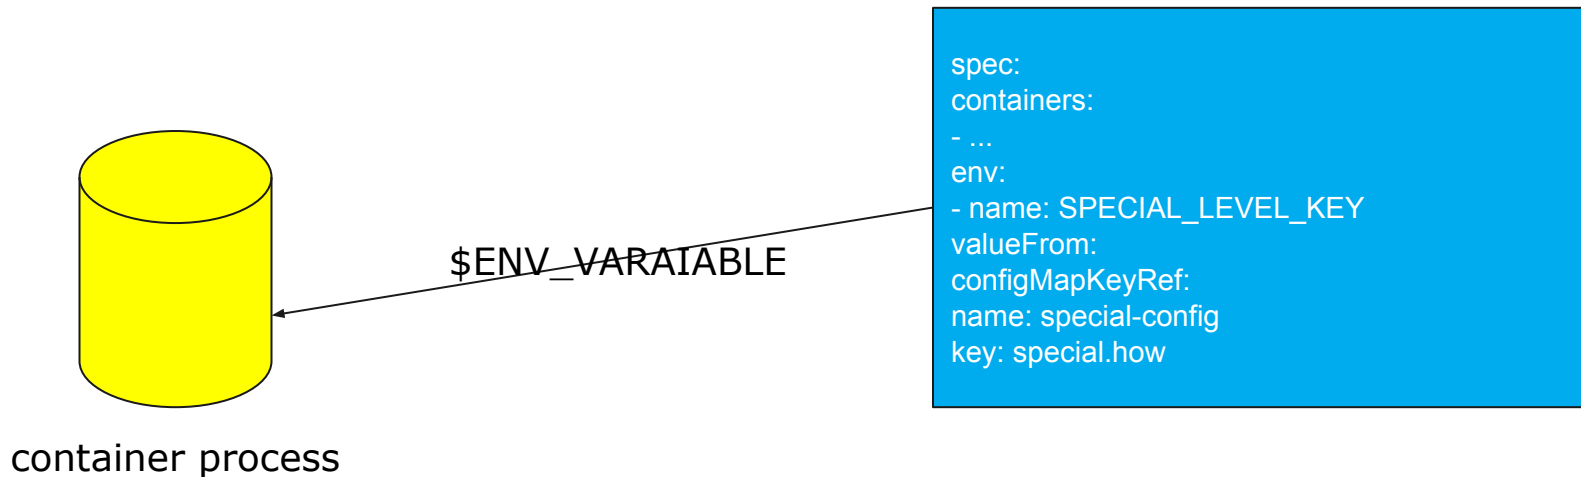
K8s Secrets

→ Ví dụ mẫu về file YAML Secrets.

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret-manifest
type: Opaque
data:
  username: YW5zaHVzY2hhdWhhbg==
  password: VGVzdGt1VybmV0ZXMxMjM0NQ==
```

Pass ConfigMap and Secrets to Containers

- User có thể chuyển Secrets và Config Map đến Container sử dụng **Environment Variables**.



Pass ConfigMap and Secrets to Containers

- Cấu hình **Mount Volume** là cách khác để chuyển dữ liệu Config và Secrets đến Container.
- Dữ liệu Config sẽ hiển thị dưới dạng Files trong hệ thống file của Container

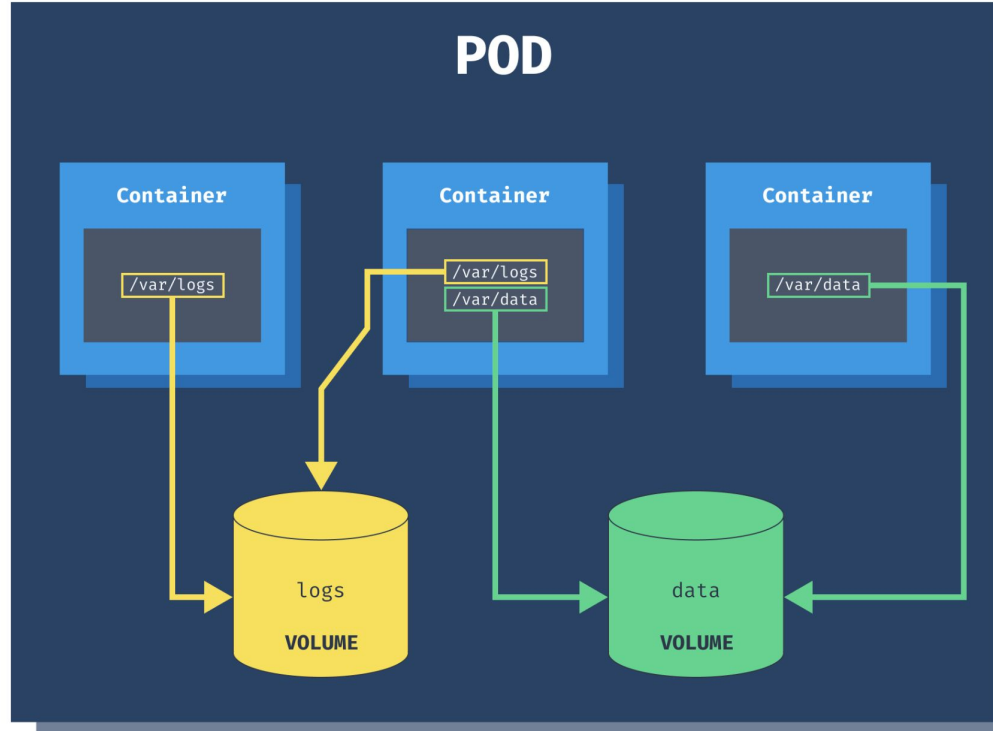
```
volumes:  
- name: config-volume  
  configMap:  
    name: special-config
```

Labs



Lab 0: Manage Application Configuration

Volume



Volume

- Volumes are special directories that are mounted in containers
- Volumes can have many different purposes:
 - share files and directories between containers running on the same machine
 - share files and directories between containers and their host
 - centralize configuration information in Kubernetes and expose it to containers
 - manage credentials and secrets and expose them securely to containers
 - store persistent data for stateful services
 - access storage systems (like Ceph, EBS, NFS, Portworx, and many others)

A simple volume example

apiVersion: v1

kind: Pod

metadata:

name: nginx-with-volume

spec:

volumes:

- name: www

containers:

- name: nginx

image: nginx

volumeMounts:

- name: www

mountPath: /usr/share/nginx/html/

A volume shared between two containers

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-with-volume
spec:
  volumes:
    - name: www
  containers:
    - name: nginx
      image: nginx
      volumeMounts:
        - name: www
          mountPath: /usr/share/nginx/html/
    - name: git
      image: alpine
      command: [ "sh", "-c", "apk add --no-cache git && git clone https://github.com/octocat/Spoon-Knife /www" ]
      volumeMounts:
        - name: www
          mountPath: /www/
      restartPolicy: OnFailure
```

Labs



Lab 1: Sharing a volume, in action

Let's try it!

1. Create the pod by applying the YAML file:
`kubectl apply -f nginx-with-volume.yaml`
2. Check the IP address that was allocated to our pod:

```
kubectl get pod nginx-with-volume -o wide  
IP=$(kubectl get pod nginx-with-volume -o json |  
jq -r .status.podIP)
```

3. Access the web server:

```
curl $IP
```

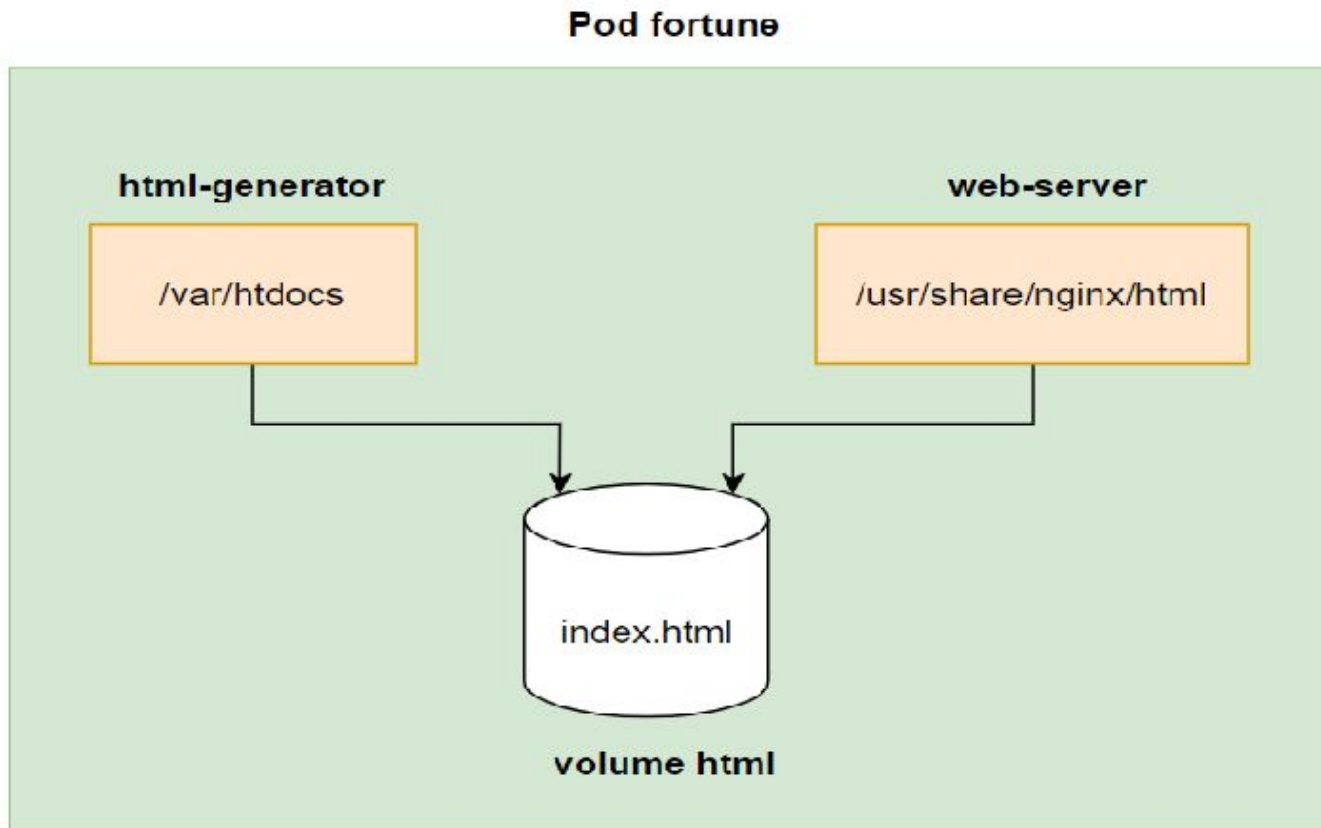
nginx-with-volume.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx-with-volume  
spec:  
  volumes:  
  - name: www  
  containers:  
  - name: nginx  
    image: nginx  
    volumeMounts:  
    - name: www  
      mountPath:  
        /usr/share/nginx/html/
```

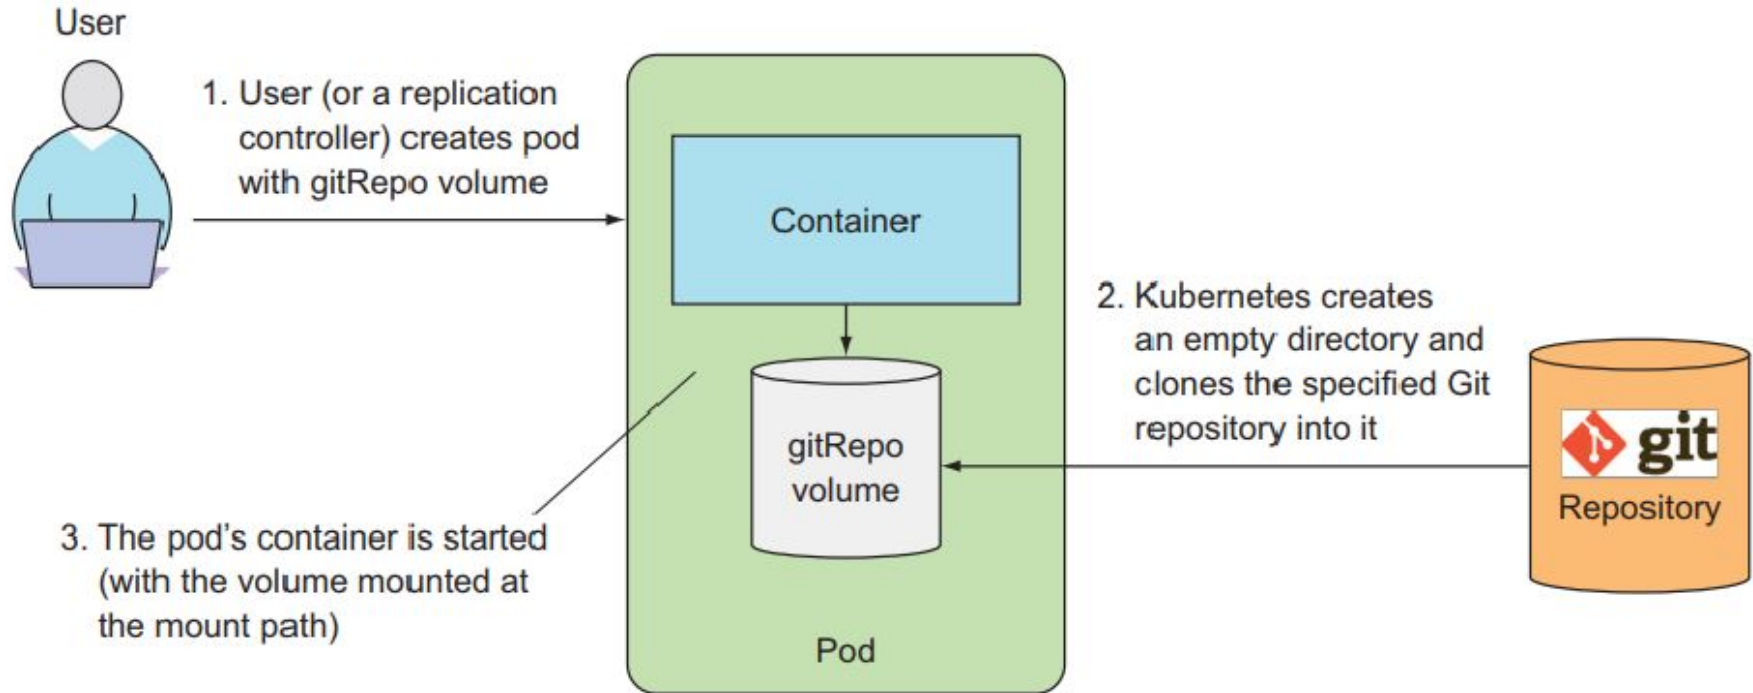
Volume Types

Temp	Local	Network
emptyDir	hostPath	GlusterFS gitRepo NFS iSCSI gcePersistentDisk AWS EBS azureDisk Fiber Channel Secret VsphereVolume

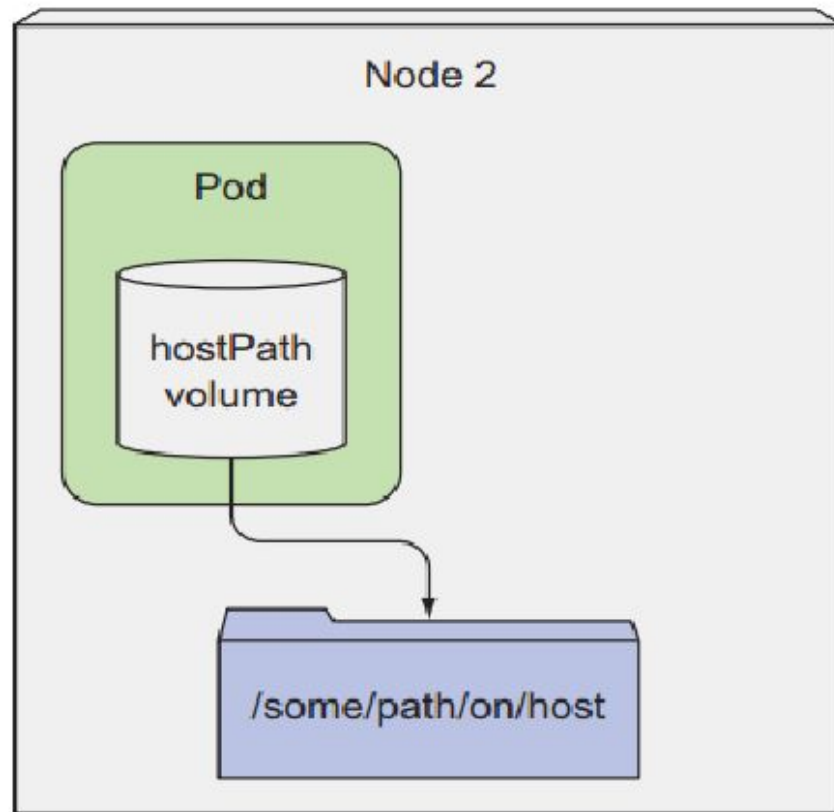
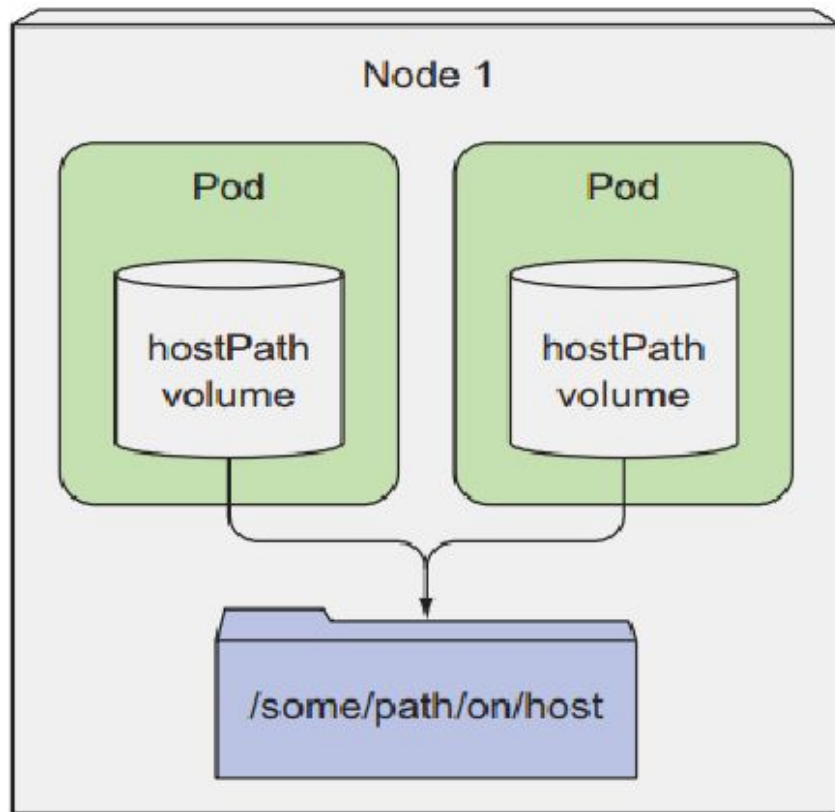
Volumes: EmptyDir



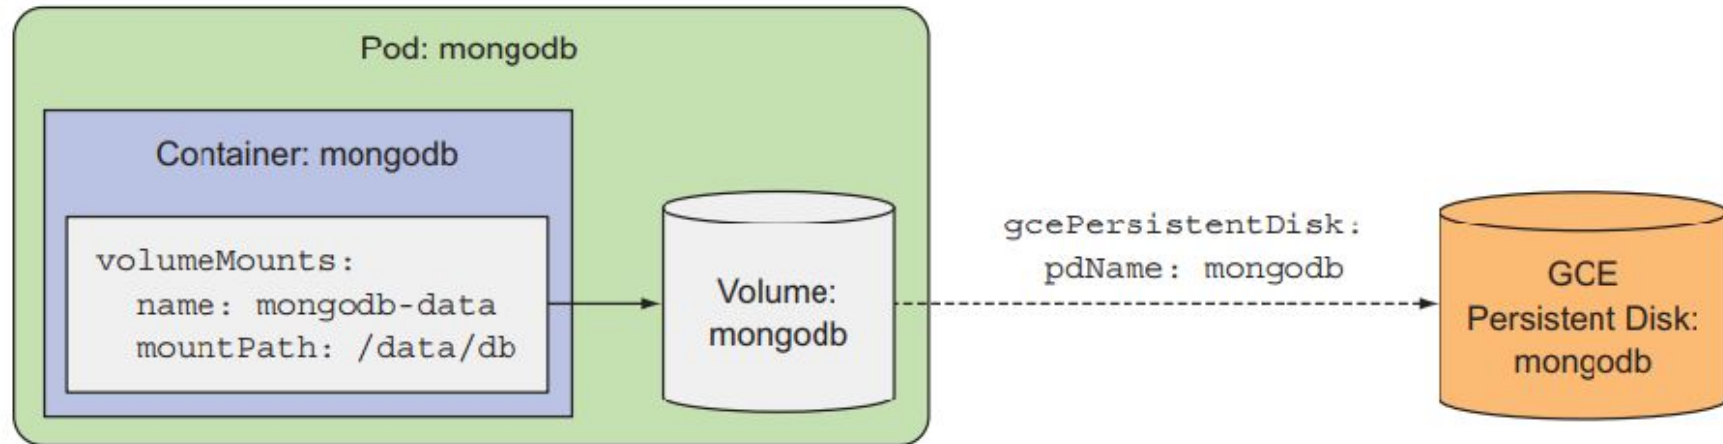
Volumes: Git repo



Volumes: hostPath



Volumes: Cloud storage



Volume lifecycle

- The lifecycle of a volume is linked to the pod's lifecycle
- This means that a volume is created when the pod is created
- This is mostly relevant for emptyDir volumes
(other volumes, like remote storage, are not "created" but rather "attached")
- A volume survives across container restarts
- A volume is destroyed (or, for remote storage, detached) when the pod is destroyed

Labs



Lab 2: Volume type: hostPath

Let's try it!

1. Create the pod by applying the YAML file:

```
kubectl apply -f
```

```
pod-volume-hostPath.yaml
```

2. Check the IP address that was allocated to our pod:

```
kubectl get pod test-pd -o wide
```

```
IP=$(kubectl get pod nginx-with-volume  
-o json | jq -r .status.podIP)
```

3. Access the web server:

```
curl $IP
```

pod-volume-hostPath.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: test-pd  
spec:  
  containers:  
    - image: k8s.gcr.io/test-webserver  
      name: test-container  
      volumeMounts:  
        - mountPath: /test-pd  
          name: test-volume  
  volumes:  
    - name: test-volume  
      hostPath:  
        # directory location on host  
        path: /data  
        # this field is optional  
        type: DirectoryOrCreate
```

Section 2:

PV, PVC and Statefulset

Volumes and Persistent Volumes

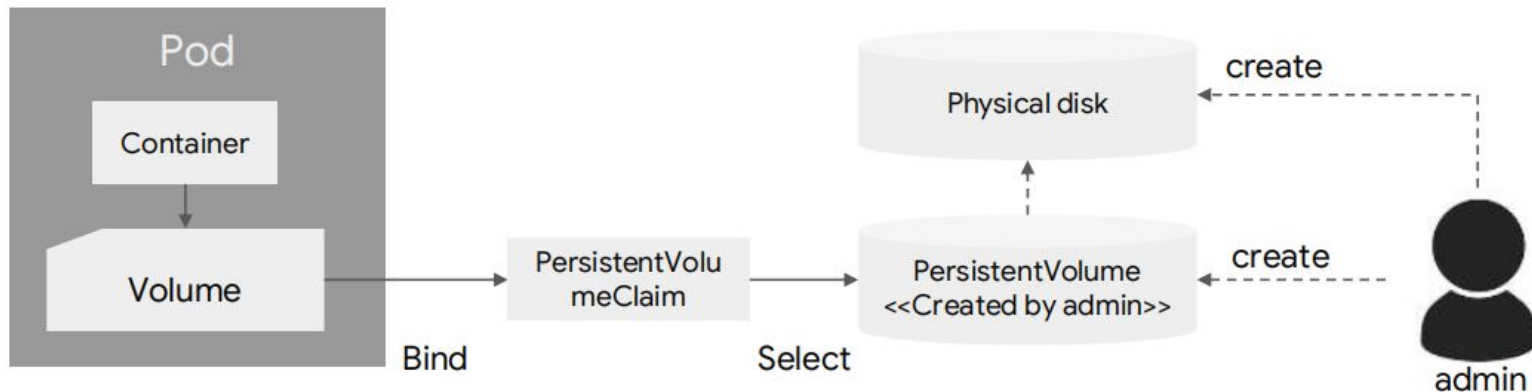
- **Volumes** are used for many purposes:
 - sharing data between containers in a pod
 - exposing configuration information and secrets to containers
 - accessing storage systems
- The last type of volumes is known as a "Persistent Volume"

Persistent Volumes types

- There are many **types of Persistent Volumes** available:
 - public cloud storage (GCEPersistentDisk, AWSElasticBlockStore, AzureDisk...)
 - private cloud storage (Cinder, VsphereVolume...)
 - traditional storage systems (NFS, iSCSI, FC...)
 - distributed storage (Ceph, Glusterfs, Portworx...)
- Using a persistent volume requires:
 - creating the volume out-of-band (outside of the Kubernetes API)
 - referencing the volume in the pod description, with all its parameters

PersistentVolume

- Life cycle is managed by k8s cluster (not pod)
- Admin can create PersistentVolume (Static provisioning) and developer just use the volume through PersistentVolumeClaim without understanding of infrastructure (It is more common to use dynamic volume provisioning instead of volume provisioning by admin)



PersistentVolume

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
  mountOptions:
    - hard
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2
```

- Capacity : Storage size
In future it will include IOPS,throughput etc
- VolumeMode (from 1.9) : Filesystem(default) or rawblock device
- Reclaim Policy
 - Retain – manual reclamation
 - Recycle – basic scrub (rm -rf /thevolume/*)
 - Delete – associated storage asset such as AWS EBS, GCE PD, Azure Disk, or OpenStack Cinder volume is deleted
Currently, only NFS and HostPath support recycling. AWS EBS, GCE PD, Azure Disk, and Cinder volumes support deletion.
- Mount Option
Additional mount options for when a Persistent Volume is mounted on a node

PersistentVolume

- AccessModes

- ReadWriteOnce (RWO)– the volume can be mounted as read-write by a single node
- ReadOnlyMany (ROX) – the volume can be mounted read-only by many nodes
- ReadWriteMany (RWX) – the volume can be mounted as read-write by many nodes

Note: A volume can only be mounted using one access mode at a time, even if it supports many. For example, a GCEPersistentDisk can be mounted as ReadWriteOnce by a single node or ReadOnlyMany by many nodes, but not at the same time.

Volume Plugin	ReadWriteOnce	ReadOnlyMany	ReadWriteMany	ReadWriteOncePod
AWSElasticBlockStore	✓	-	-	-
AzureFile	✓	✓	✓	-
AzureDisk	✓	-	-	-
CephFS	✓	✓	✓	-
Cinder	✓	-	(if multi-attach volumes are available)	-
CSI	depends on the driver	depends on the driver	depends on the driver	depends on the driver
FC	✓	✓	-	-
FlexVolume	✓	✓	depends on the driver	-
GCEPersistentDisk	✓	✓	-	-
Glusterfs	✓	✓	✓	-
HostPath	✓	-	-	-
iSCSI	✓	✓	-	-
NFS	✓	✓	✓	-
RBD	✓	✓	-	-
VsphereVolume	✓	-	- (works when Pods are collocated)	-
PortworxVolume	✓	-	✓	-

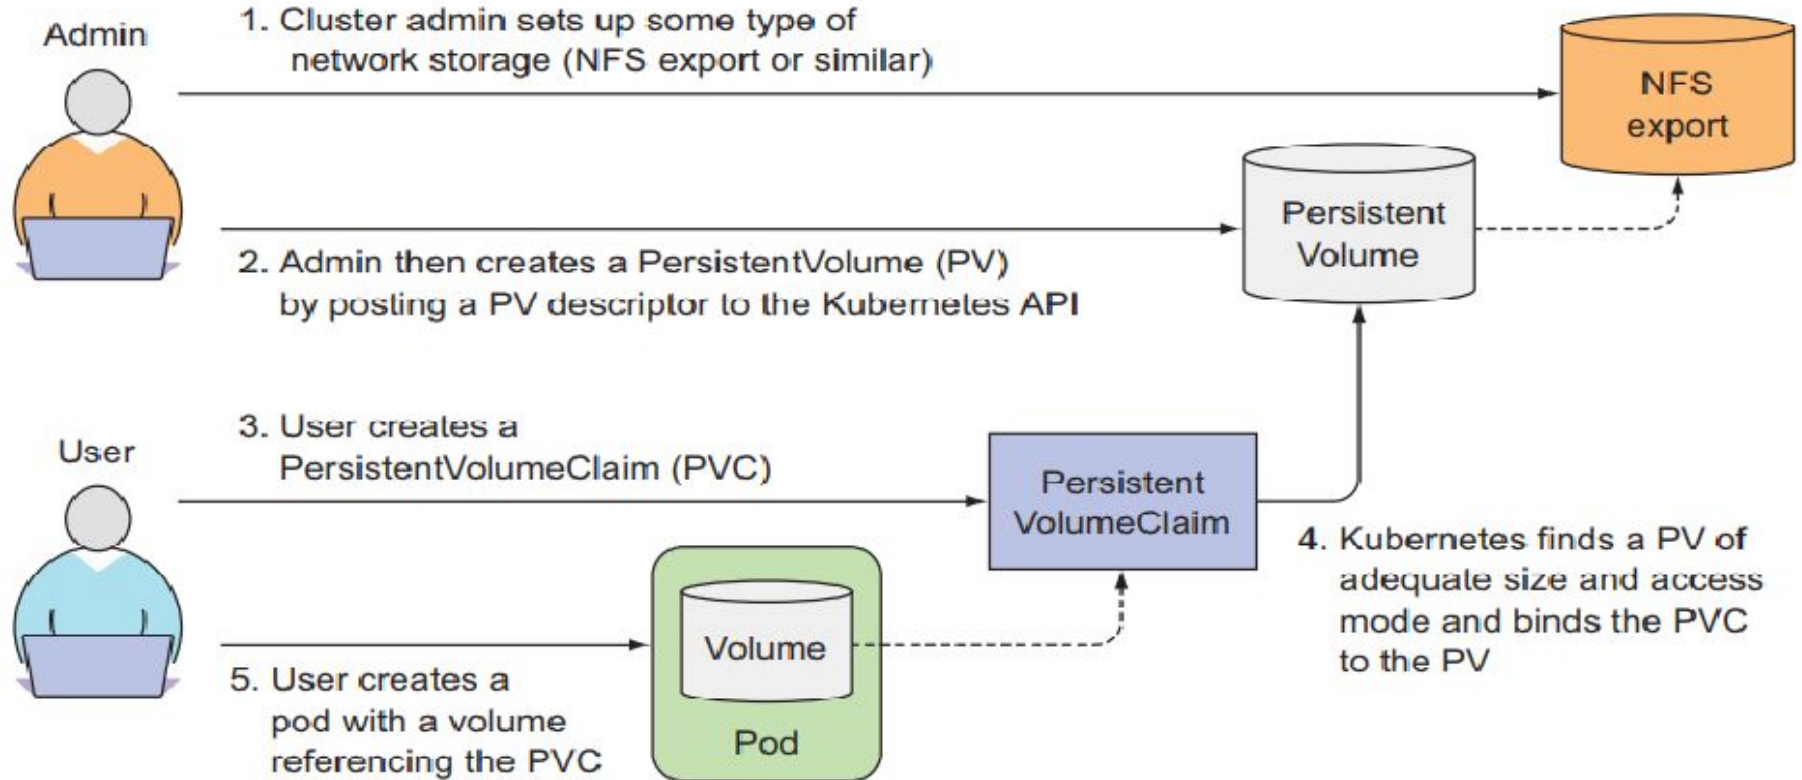
PersistentVolume Phase

- Available - a free resource that yet is not bound to a claim
- Bound - the volume is bound to a claim
- Released - the claim has been deleted but the resource is not yet reclaimed by the cluster
- Failed - the volume has failed its automatic reclamation

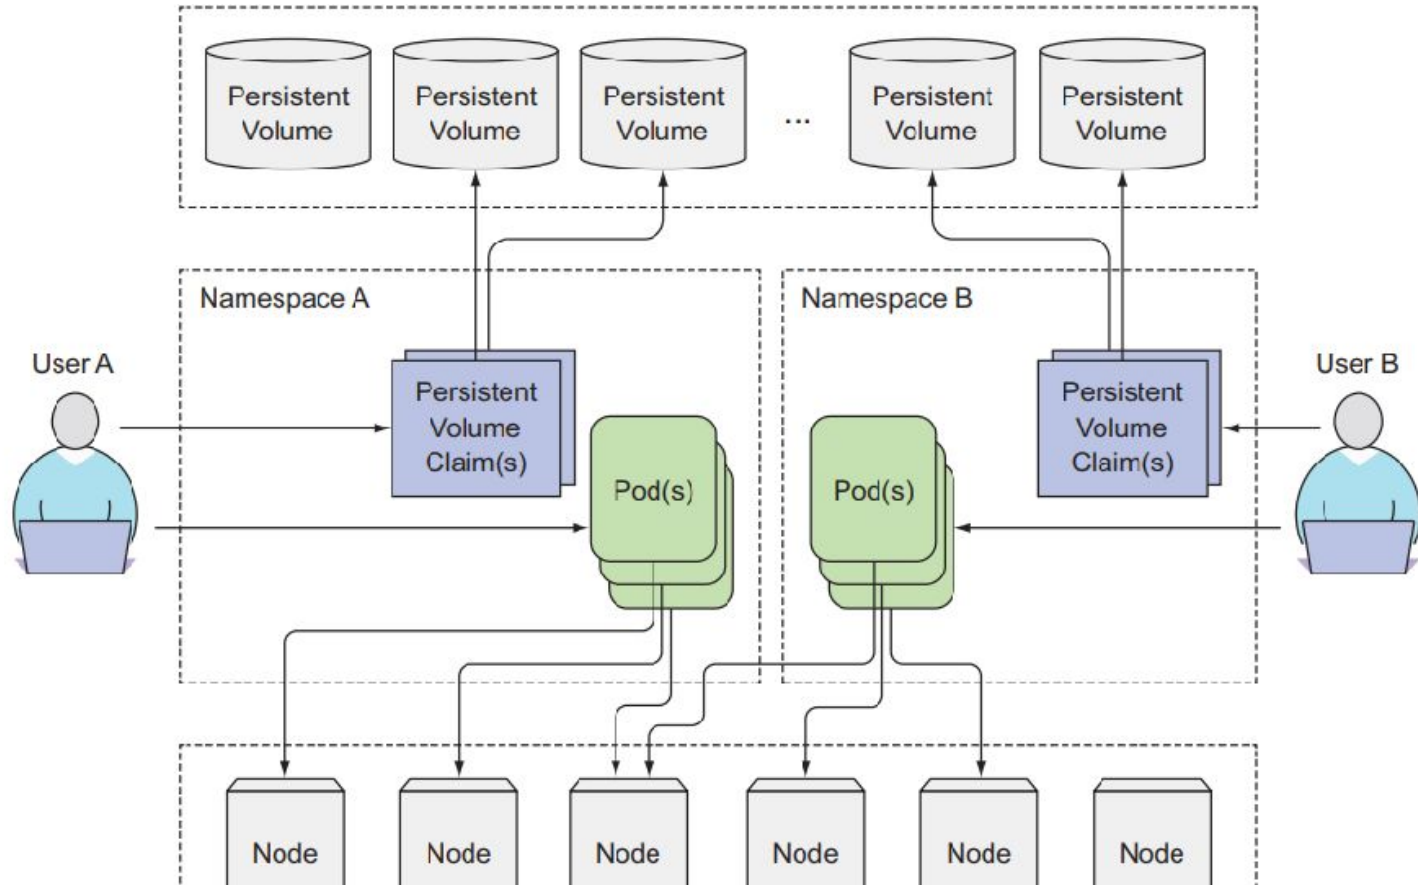
Lifecycle of volume and claim

1. Provisioning
 - Static
 - Dynamic : Dynamically create volume with storage class.
2. Binding
 - Bind PV to PVC
3. Using
 - Bind PVC to Pod and starts to use
4. Reclaiming

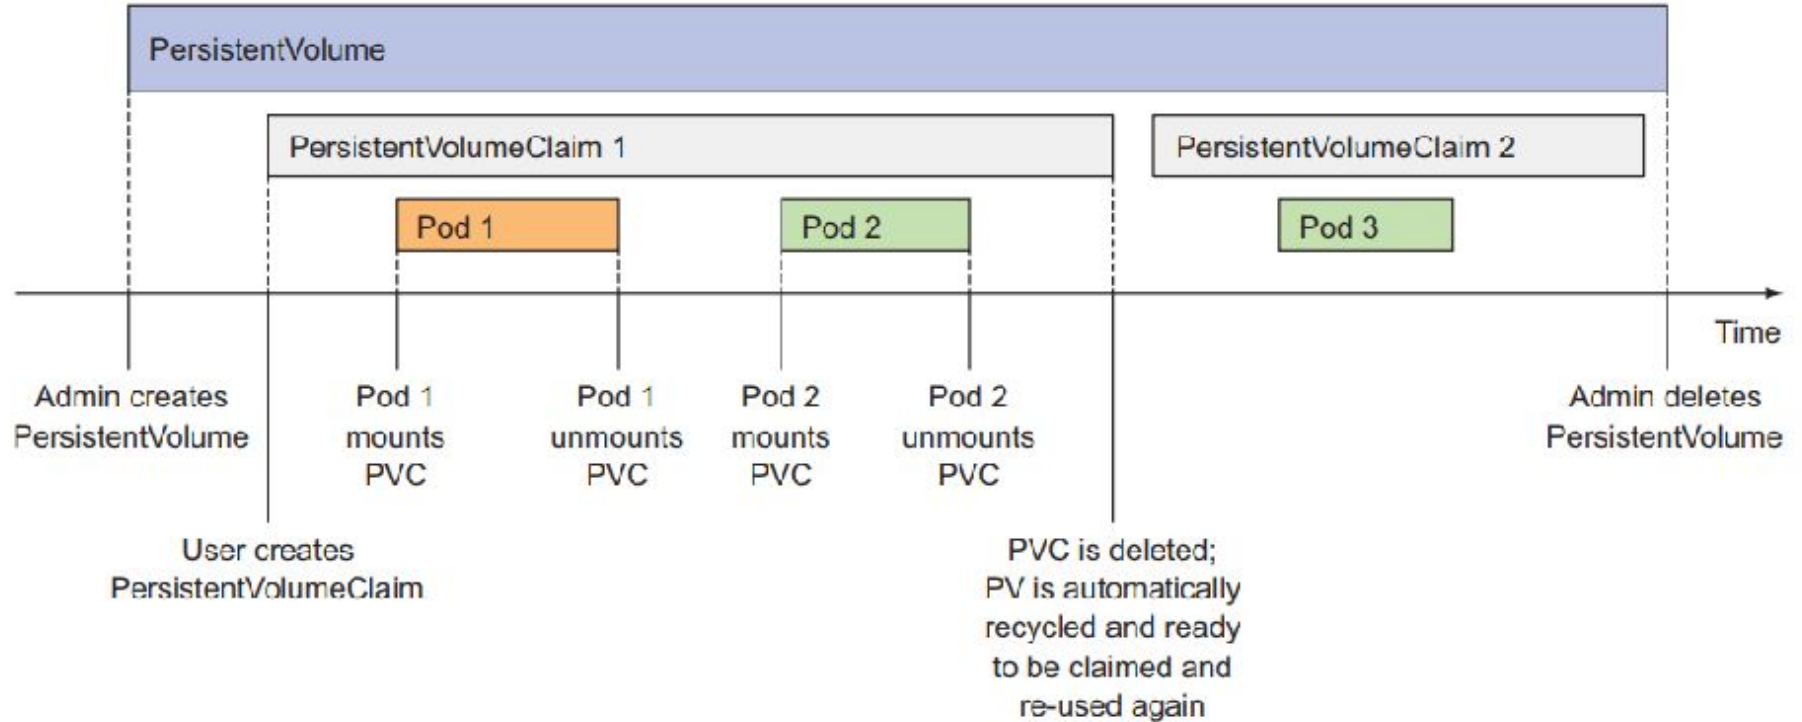
PersistentVolumeClaim (PVC)



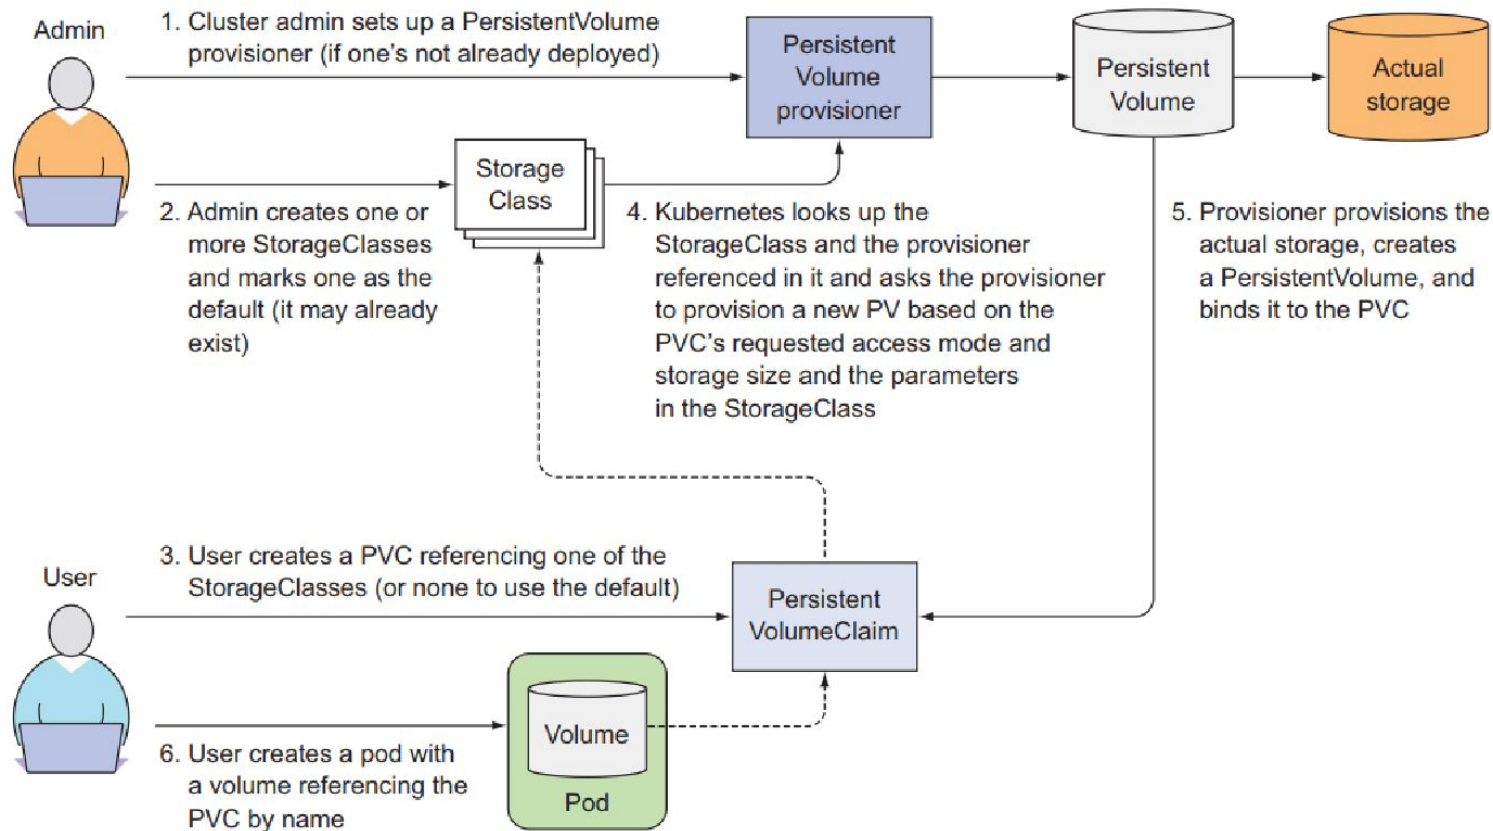
PersistentVolumeClaim (PVC)



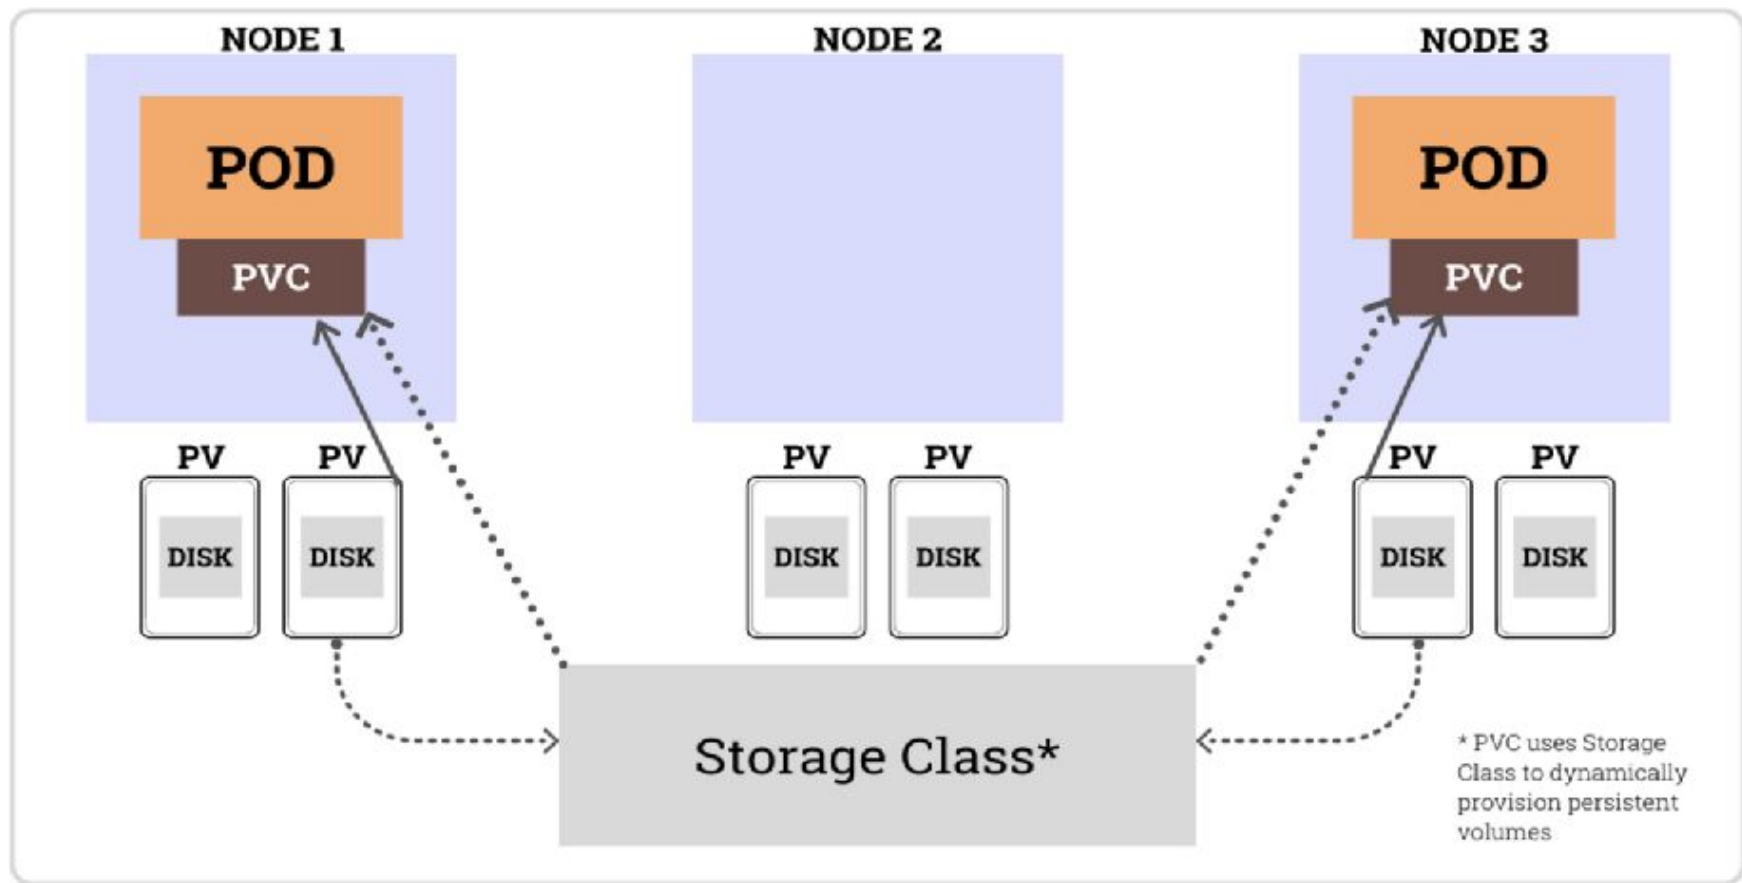
Recycling PersistentVolumes



StorageClass



StorageClass



Labs



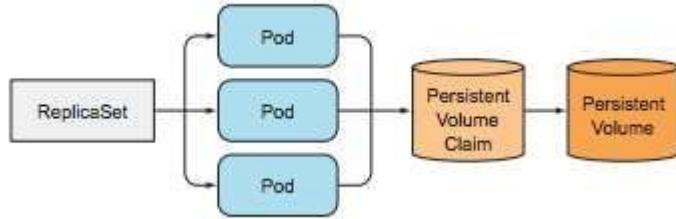
Lab 3: Using PV, PVC

Shortcomings of Persistent Volumes


- Their lifecycle (creation, deletion...) is managed outside of the Kubernetes API (we can't just use `kubectl apply/create/delete/...` to manage them)
- If a Deployment uses a volume, all replicas end up using the same volume
- That volume must then support concurrent access
 - some volumes do (e.g. NFS servers support multiple read/write access)
 - some volumes support concurrent reads
 - some volumes support concurrent access for colocated pods
- What we really need is a way for each replica to have its own volume

Replicating stateful pod

- With ReplicaSet



- Option : Create RS per Pod



Statefulset

- Stateful sets are a type of resource in the Kubernetes API (like pods, deployments, services...)
- They offer mechanisms to deploy scaled stateful applications
- At a first glance, they look like *deployments*:
 - a stateful set defines a pod spec and a number of replicas R
 - it will make sure that R copies of the pod are running
 - that number can be changed while the stateful set is running
 - updating the pod spec will cause a rolling update to happen
- But they also have some significant differences

Stateful sets unique features

- Pods in a stateful set are numbered (from 0 to $R-1$) and ordered
- They are started and updated in order (from 0 to $R-1$)
- A pod is started (or updated) only when the previous one is ready
- They are stopped in reverse order (from $R-1$ to 0)
- Each pod know its identity (i.e. which number it is in the set)
- Each pod can discover the IP address of the others easily
- The pods can have persistent volumes attached to them

Statefulset

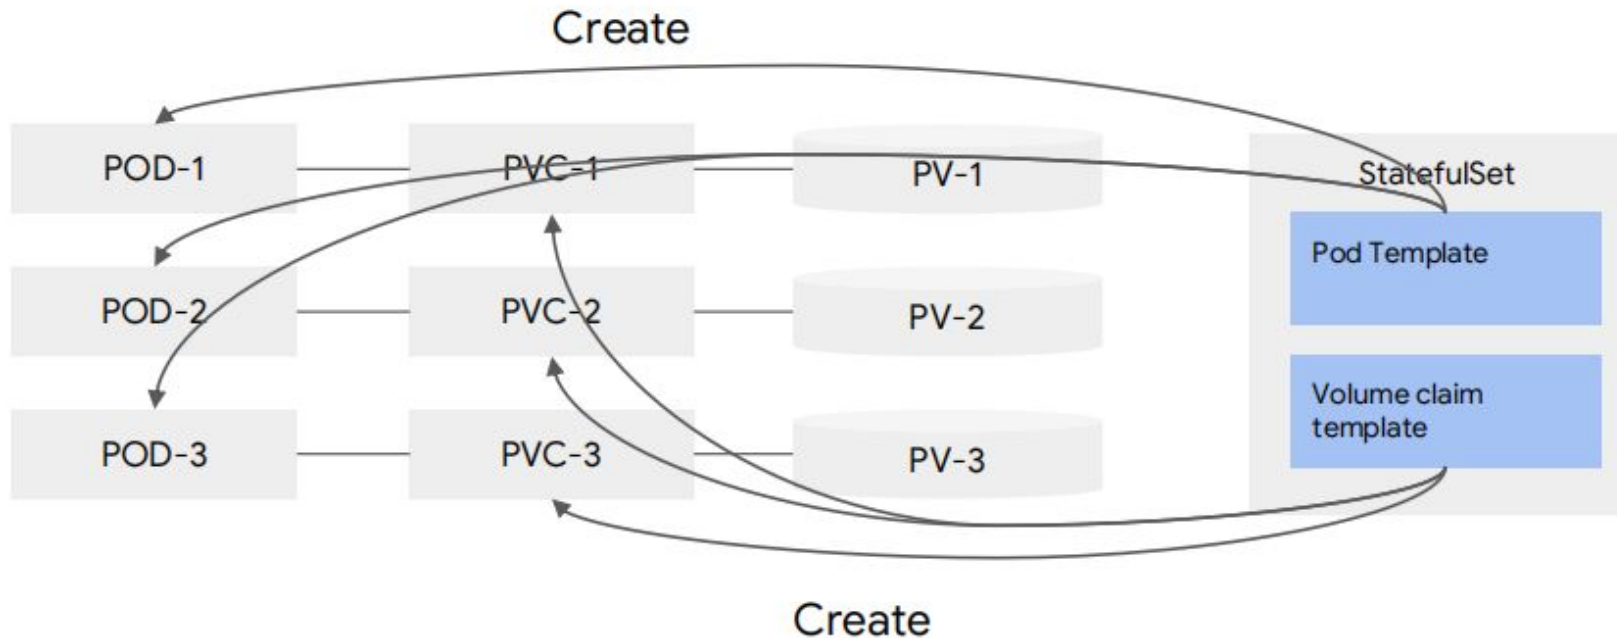
- To support stateful application (like database)
- It is supported from k8s 1.9(GA)
- Pod Naming : `${Stateful set name}-${ordinary index}`
It is not that the pods having a predictable name and hostname.
It is better to use service for naming.
- When Pod (under statefulset) is restarted (by crash), Pod name will not be changed. Not like RS, the Pod name is changed to new one.

StatefulSet specification example

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  selector:
    matchLabels:
      app: nginx # has to match .spec.template.metadata.labels
  serviceName: "nginx"
  replicas: 3 # by default is 1
  template:
    metadata:
      labels:
        app: nginx # has to match .spec.selector.matchLabels
    spec:
      terminationGracePeriodSeconds: 10
      containers:
        - name: nginx
          image: k8s.gcr.io/nginx-slim:0.8
          ports:
            - containerPort: 80
              name: web
          volumeMounts:
            - name: www
              mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
    - metadata:
        name: www
      spec:
        accessModes: [ "ReadWriteOnce" ]
        storageClassName: "my-storage-class"
      resources:
        requests:
          storage: 1Gi
```

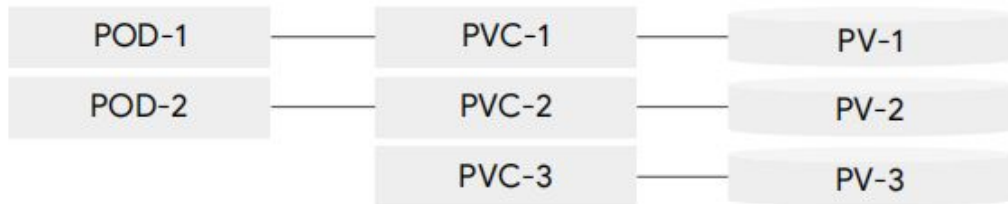
Pod to PV mapping

- Each pod will get its own PVC and PV
- The PVC will be created “Volume claim template” in StatefulSet specification



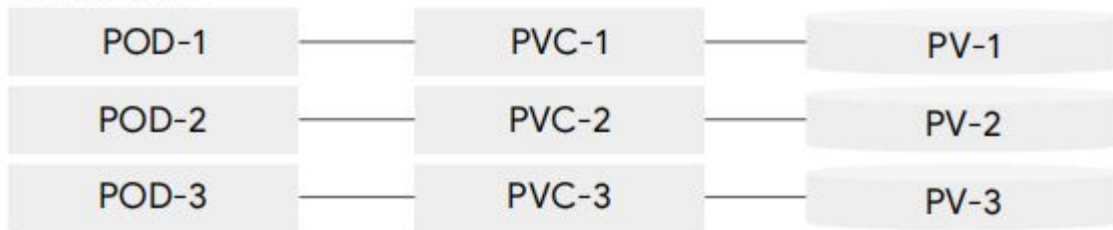
Scale in & out

- Scale in



POD-3 is removed by scale in , But PVC-3 and PV-3 are not deleted

- Scale out



POD-3 is added by scale out. It is attached previous PVC-3 and PV-3

It enables POD connect to same PVC and PV and it POD to retain same stateful information in disk

Labs



Lab 4: StatefulSets



Section 3:

Summary

Mục tiêu

Kết thúc bài học, học viên cần đạt được các kỹ năng sau:

- Có thể hiểu và tạo được Volume trong Kubernetes.
- Nắm được các đối tượng để khởi tạo Volume như StorageClass, PersistentVolume và PersistentVolumeClaim.
- Có khả năng triển khai ứng dụng dưới dạng Statefulsets.

Tài liệu tham khảo

Thank you